

Weekly Assignment

Aggregate and Atomicity

1. What is an aggregate function in SQL? Give an example?

An aggregate function in SQL performs a calculation on a set of values and returns a single value. Common aggregate functions include `SUM()`, `AVG()`, `COUNT()`, `MAX()`, and `MIN()`.

Example:

```
SELECT AVG(salary) AS average_salary FROM employees;
```

2. How can you use the GROUP BY clause in combination with aggregate functions?

The GROUP BY clause groups rows that have the same values in specified columns into summary rows. It is often used with aggregate functions to perform calculations on each group.

Example:

```
SELECT department, AVG(salary) AS average_salary FROM employees GROUP BY department;
```

3. Describe a scenario where atomicity is crucial for database operations.

Atomicity is a fundamental property of database transactions that ensures operations are completed entirely or not at all, maintaining the database's integrity. A scenario where atomicity is crucial involves financial transactions, such as transferring money between two bank accounts.

Importance of Atomicity:

1. **Consistency of Data:** Ensures data remains accurate and consistent.
2. **Avoiding Partial Updates:** Prevents partial updates that could lead to data anomalies.
3. **Maintaining Integrity:** Ensures all operations within a transaction are completed or none at all.

OLAP and OLTP

1. Mention any 2 differences between OLAP and OLTP

1. Purpose and Use Cases

- OLAP (Online Analytical Processing):
 - Purpose: Designed for complex querying and data analysis. Used for data mining, trend analysis, and business intelligence.
 - Use Cases: Generating reports on sales performance, analyzing market trends, creating dashboards for executive summaries.
- OLTP (Online Transaction Processing):
 - Purpose: Designed for managing and processing real-time transactional data. Optimized for handling a high volume of short online transactions.
 - Use Cases: Order processing systems, banking transactions, customer relationship management (CRM) systems.

2. Data Structure and Query Types

- OLAP:
 - Data Structure: Often uses multi-dimensional data models (e.g., star schema, snowflake schema) optimized for complex queries and aggregations.
 - Query Types: Generally complex, involving aggregations, calculations, and multidimensional analysis (e.g., "What is the total sales revenue for each region over the last five years?").
- OLTP:
 - Data Structure: Uses normalized relational data models to reduce redundancy and ensure data integrity.
 - Query Types: Usually simple, involving retrieving or updating individual records (e.g., "What is the balance of account number 123456?").

2. How do you optimize an OLTP database for better performance?

1. Use Indexes Effectively:
 - Create Indexes: On frequently queried columns, including primary and foreign keys.
 - Choose Index Types: Use B-Tree for general cases, hash indexes for equality checks, and composite indexes for multi-column queries.
 - Maintain Indexes: Regularly rebuild and monitor their usage to avoid fragmentation.
2. Optimize Queries:
 - Write Efficient Queries: Specify columns, use **WHERE** clauses, and minimize joins.
 - Avoid Complex Joins/Subqueries: Simplify queries for faster execution.
3. Optimize Database Design:
 - Normalize Data: Reduce redundancy and improve integrity.
 - Consider Controlled Denormalization: For specific performance gains.
4. Optimize Transaction Management:
 - Minimize Transaction Scope: Keep transactions short.
 - Select Appropriate Isolation Levels: Balance consistency with performance.

5. Optimize Hardware and Configuration:
 - Improve Disk I/O: Use SSDs for faster access.
 - Tune Parameters: Adjust buffer pool, cache, and log file settings.
6. Regular Maintenance:
 - Update Statistics: Keep database statistics current.
 - Monitor Performance: Use tools to track and address bottlenecks.

Data Encryption and Storage

1. What are the different types of data encryption available in MSSQL?

- Transparent Data Encryption (TDE)
- Column-Level Encryption
- Always Encrypted
- Backup Encryption
- Dynamic Data Masking

SQL, NoSQL, Applications, Embedded

1. What is the main difference between SQL and NoSQL databases?

Aspect	SQL Databases	NoSQL Databases
Data Model	Relational, structured in tables	Non-relational, various formats (key-value, documents, graphs, wide-column)
Schema	Fixed schema	Flexible schema
Examples	MySQL, PostgreSQL, SQL Server, Oracle	MongoDB (document-based), Redis (key-value store), Neo4j (graph database), Cassandra (wide-column store)

DDL

1. How do you create a new schema in MSSQL?

A schema is a blueprint that outlines the tables, views, indexes, relationships, and constraints within a database.

Syntax:

```
CREATE SCHEMA schema_name;
```

2. Describe the process of altering an existing table

Add a Column:

```
ALTER TABLE table_name ADD new_column_name data_type;
```

Modify a Column:

```
ALTER TABLE table_name ALTER COLUMN column_name new_data_type;
```

Drop a Column:

```
ALTER TABLE table_name DROP COLUMN column_name;
```

Add a Constraint:

```
ALTER TABLE table_name ADD CONSTRAINT constraint_name constraint_type  
(column_name);
```

Drop a Constraint:

```
ALTER TABLE table_name DROP CONSTRAINT constraint_name;
```

3. What is the difference between a VIEW and a TABLE in MSSQL?

Aspect	Table	View
Definition	A physical structure that stores data in rows and columns	A virtual table based on a SELECT query
Data Storage	Stores data persistently on disk	Does not store data; provides a dynamic view of data
Schema	Fixed schema with defined columns and constraints	Defined by a SELECT query; schema can change based on query
Data Modification	Directly insert, update, or delete data	Data modification depends on the view's complexity and underlying tables
Indexes	Can have indexes to improve query performance	Does not have indexes; indexed views can be created for performance
Usage	Used for storing actual data	Used for presenting data, simplifying complex queries, and providing security

4. Explain how to create and manage indexes in a table.

Creating a Simple Index A simple index is created on a single column and is used to speed up searches on that column.

Syntax:

```
CREATE INDEX index_name ON table_name (column_name);
```

Managing Indexes

Viewing Existing Indexes To view existing indexes on a table, you can query the `sys.indexes` catalog view or use SQL Server Management Studio (SSMS).

```
SELECT * FROM sys.indexes WHERE object_id = OBJECT_ID('table_name');
```

Rebuilding an Index Rebuilding an index can help defragment and optimize its performance. This is typically done during maintenance operations.

```
ALTER INDEX index_name ON table_name REBUILD;
```

DML

1. What are the most commonly used DML commands?

- **SELECT:** Retrieves data.
- **INSERT:** Adds new data.
- **UPDATE:** Modifies existing data.
- **DELETE:** Removes data.
- **MERGE:** Combines insert, update, and delete operations based on a condition.

2. How do you retrieve data from multiple tables using a JOIN?

One of the most common approaches to retrieve data from multiple tables in SQL is by utilizing `JOIN` clauses to combine data from different tables based on specified conditions.

Syntax:

```
SELECT t1.column1, t2.column2  
  
FROM table1 t1  
  
JOIN table2 t2 ON t1.id = t2.id;
```

3. Explain how relational algebra is used in SQL queries.

Relational algebra provides a theoretical foundation for SQL operations by defining how data can be manipulated and combined. SQL queries implement these operations to retrieve and manipulate data from relational databases:

- **Selection (σ)** is equivalent to `WHERE`.

- **Projection (π)** is equivalent to **SELECT**.
- **Union (\cup)** is equivalent to **UNION**.
- **Intersection (\cap)** is equivalent to **INTERSECT**.
- **Difference ($-$)** is equivalent to **EXCEPT**.
- **Cartesian Product (\times)** is equivalent to **CROSS JOIN**.
- **Join (\bowtie)** is equivalent to various types of **JOIN** operations.
- **Division (\div)** can be more complex and often requires subqueries.

4. What are the implications of using complex queries in terms of performance and maintainability?

Implications of Complex Queries:

1. Performance:
 - Increased Processing Time: Complex queries often require more processing power and time, especially with large datasets.
 - Resource Utilization: They can consume significant database resources (CPU, memory, I/O), potentially impacting other operations.
 - Optimization Challenges: Complex queries may require careful indexing, query optimization, and sometimes rewriting for performance.
2. Maintainability:
 - Readability: Complex queries can be difficult to read and understand, making maintenance challenging for developers.
 - Debugging: Identifying and fixing issues in complex queries can be more time-consuming.
 - Modification: Changes to business logic or requirements may require substantial query modifications, increasing the risk of introducing errors.
 - Documentation: Proper documentation is essential to ensure that the purpose and functionality of complex queries are clear to other developers.

Indexes & Constraints

1. **What are indexes and why are they used?**
 - **Indexes** are database objects that improve the speed of data retrieval operations on a table at the cost of additional space and potentially slower data modification operations. They work similarly to an index in a book, allowing the database to quickly locate specific rows in a table without scanning the entire table.
2. **What Are Indexes?** An index is a data structure that provides a fast way to look up and retrieve rows from a database table based on the values in one or more columns. In most database systems, indexes are implemented using data structures like B-trees or hash tables.

Uses:

 - **Faster Queries:** Speed up data searches and retrieval by allowing quick access to rows without scanning the entire table.

- **Efficient Sorting:** Help with fast sorting and ordering of query results.
- **Optimized Joins:** Enhance the performance of join operations by quickly matching rows between tables.
- **Uniqueness Enforcement:** Unique indexes ensure that values in specified columns are unique, maintaining data integrity.
- **Better Aggregations:** Improve the performance of aggregate functions like COUNT, SUM, and AVG.

3. **How do you create a unique constraint on a table column?**

- To create a unique constraint on a table column in SQL, you use the **UNIQUE** keyword. A unique constraint ensures that all values in a column or a combination of columns are distinct across the table, which helps maintain data integrity by preventing duplicate entries.

Example:

sql

```
CREATE TABLE employees (
    employee_id INT PRIMARY KEY,
    email VARCHAR(255) UNIQUE,
    name VARCHAR(100)
);
```

- 4.
5. **Explain the difference between clustered and non-clustered indexes.**

Feature	Clustered Index	Non-Clustered Index
Data Storage	Alters the physical order of the table data	Maintains a separate structure with pointers to the data
Number per Table	One per table	Multiple per table

Performance for Queries	Optimal for queries that benefit from sorted data (range queries, ordered results)	Efficient for lookups, searches, and queries on non-primary key columns
Impact on Data Modification	Can impact insert/update performance due to physical reordering	May slow down insert/update operations due to index maintenance

4. **How would you optimize index usage in a highly transactional database?**
 - **Optimizing index usage** in a highly transactional database involves carefully balancing query performance with the overhead introduced by maintaining indexes.
 - Analyzing query patterns to design effective indexes.
 - Monitoring and maintaining indexes to avoid performance degradation.
 - Balancing index creation with data modification performance.
 - Using partitioning or sharding for very large datasets.
 - Optimizing queries to ensure effective index usage.

Joins

1. **What are the different types of joins available in MSSQL?**
 - **Inner Join:** Returns rows with matching values in both tables.
 - **Left Join:** Returns all rows from the left table and matched rows from the right table (NULLs if no match).
 - **Right Join:** Returns all rows from the right table and matched rows from the left table (NULLs if no match).
 - **Full Join:** Returns all rows with matches from either table (NULLs where no match).
 - **Cross Join:** Returns the Cartesian product of the tables.
 - **Self Join:** Joins the table with itself to compare rows within the same table.

Provide an example of a LEFT JOIN query.

sql

```
SELECT employees.name, departments.department_name
FROM employees
LEFT JOIN departments ON employees.department_id =
departments.department_id;
```

2.

3. **Explain the concept of a self-join and when it might be used.**

- **Definition:** Joins a table with itself to compare rows within the same table.
- **Usage:** Useful for hierarchical data or when you need to compare rows within the same table.

Syntax:

sql

```
SELECT a.columns, b.columns
```

```
FROM table a
```

```
INNER JOIN table b ON a.column = b.column;
```

○

4. **How do you perform a full outer join and what is its significance?**

- **Definition:** Returns all rows when there is a match in either left or right table. If there is no match, NULL values are returned for columns from the table without a match.
- **Usage:** Useful when you need all records from both tables, with NULLs for non-matching rows.

Syntax:

sql

```
SELECT columns
```

```
FROM table1
```

```
FULL JOIN table2 ON table1.column = table2.column;
```

○

Alias

1. **What is an alias in SQL and how is it used?**

- In SQL, an **alias** is a temporary name assigned to a table or a column within a query. Aliases are used to simplify complex queries, make the query results more readable, or provide a more meaningful name for columns and tables. Aliases are especially useful in scenarios involving multiple tables or complex calculations.

Using Aliases for Columns:

sql

```
SELECT column_name AS alias_name FROM table_name;
```

Using Aliases for Tables:

sql

```
SELECT column_name FROM table_name AS alias_name;
```

2.

Give an example of using table aliases in a query.

sql

```
SELECT e.first_name, e.last_name
```

```
FROM employees AS e;
```

3.

4. How do you use column aliases in conjunction with aggregate functions?

- Using column aliases in conjunction with aggregate functions in SQL allows you to provide meaningful names to the results of aggregate calculations, making your query results easier to understand. Aggregate functions perform a calculation on a set of values and return a single value. Common aggregate functions include **SUM**, **COUNT**, **AVG**, **MIN**, and **MAX**.

Syntax:

sql

```
SELECT aggregate_function(column_name) AS alias_name
```

```
FROM table_name
```

```
GROUP BY column_name;
```

5.

6. Explain the benefits of using aliases in complex queries.

- **Improved Readability:**
 - **Simplify References:** Aliases shorten table and column names, making queries easier to read and write, especially when dealing with long or complex names.
- **Manage Complex Queries:**
 - **Multiple Joins:** In queries involving multiple joins, aliases help differentiate between columns from different tables.
- **Enhance Query Maintainability:**

- **Easier Updates:** Queries with aliases are easier to modify and maintain. Changing a table or column name requires only a single update in the alias definition rather than throughout the entire query.
 - **Consistent Naming:** Aliases ensure consistent naming conventions, which helps avoid confusion when querying complex data structures.
- **Resolve Naming Conflicts:**
 - **Avoid Ambiguities:** Aliases resolve conflicts when different tables have columns with the same names, such as during joins.
- **Improve Performance:**
 - **Optimized Execution:** While aliases themselves don't directly improve performance, they can make the query more manageable and help in optimizing the execution plan by clearly defining table and column relationships.
- **Facilitate Query Design and Debugging:**
 - **Design Queries:** Aliases help in designing queries by simplifying the logical structure, especially in cases involving nested queries or multiple layers of data retrieval.
 - **Debugging:** Aliases make it easier to identify and troubleshoot issues by providing clear, descriptive names for each part of the query.

Joins vs Subqueries

1. What is the difference between joins and subqueries?

Feature	Joins	Subqueries
Performance	Typically more efficient for combining large datasets due to optimized join algorithms.	Can be less efficient, especially if used in the WHERE clause, as they may be executed multiple times for each row of the outer query.
Complexity	Generally clearer when combining data from multiple tables and are more intuitive for many users.	Useful for complex conditions but might be harder to understand and manage, especially in deeply nested scenarios.
Use Cases	Preferred for straightforward data retrieval from multiple tables where relationships are defined.	Useful for filtering results, performing operations on intermediate results, or

when a query needs to use results from another query.

Readability	Tend to be more readable for many users, especially when working with relational data.	Can be less readable, particularly when they are deeply nested.
--------------------	--	---

2.

When would you prefer a subquery over a join?

- Filtering based on aggregated or calculated results.
- Working with complex conditions or comparisons.
- Avoiding Cartesian products.
- Using scalar results for comparison.
- Performing in situ calculations.
- Handling cases where joins are inefficient.
- Keeping queries simpler and more manageable.

3. **Explain how correlated subqueries work with an example.**

- A correlated subquery is a type of subquery that references columns from the outer query. Unlike a regular subquery, which is self-contained and can be executed independently, a correlated subquery is executed once for each row processed by the outer query. This means the inner query (the subquery) depends on the outer query for its values, and its results vary based on the values of the current row in the outer query.

Example:

sql

```
SELECT e1.employee_id, e1.name, e1.salary
```

```
FROM employees e1
```

```
WHERE e1.salary > (SELECT AVG(e2.salary) FROM employees e2 WHERE  
e2.department_id = e1.department_id)
```

Multiple Table Queries

1.How do you retrieve data from multiple tables using joins?

sql

```
SELECT columns
```

```
FROM table1
```

```
JOIN table2 ON table1.column = table2.column;
```

2. Provide an example of a complex query involving multiple types of joins.

```
SELECT e.name, d.department_name, p.project_name
```

```
FROM employees e
```

```
LEFT JOIN departments d ON e.department_id = d.department_id
```

```
INNER JOIN projects p ON e.project_id = p.project_id;
```

3. Explain the use of subqueries in the FROM clause with an example.

- Subqueries in the **FROM** clause, often referred to as derived tables or inline views, allow you to define a temporary result set that can be referenced within the main query. This is useful for breaking down complex queries into more manageable parts or for performing intermediate calculations that need to be used in the main query.

Example:

```
SELECT avg_dept_salary.department_id, avg_dept_salary.avg_salary
```

```
FROM (SELECT department_id, AVG(salary) AS avg_salary
```

```
      FROM employees
```

```
      GROUP BY department_id) AS avg_dept_salary;
```

4) What is a cross join and when might you use it?

- **Definition:** A cross join returns the Cartesian product of the two tables, meaning it returns all possible combinations of rows from both tables.
- **Usage:** Use a cross join when you need to generate combinations of rows from two tables. This might be useful in scenarios like generating test data, creating combinations for analysis, or when combining data without a direct relationship.

Syntax:

```
SELECT *
```

```
FROM table1
```

```
CROSS JOIN table2;
```

Self Joins

1. What is a self-join and how is it useful?

- A self-join is a join where a table is joined with itself. This is useful for comparing rows within the same table, such as finding relationships or hierarchies.

Example:

sql

```
SELECT a.employee_id, a.name, b.name AS manager_name  
  
FROM employees a  
  
JOIN employees b ON a.manager_id = b.employee_id;
```

2. Provide an example of a self-join query.

sql

```
SELECT a.employee_id, a.name, b.name AS manager_name  
  
FROM employees a  
  
JOIN employees b ON a.manager_id = b.employee_id;
```

3. Explain the concept of a hierarchical query using a self-join.

- **Definition:** A hierarchical query retrieves data that is organized in a tree-like structure, such as organizational charts, family trees, or file systems.

Example:

sql

```
SELECT employee_id, name, manager_id  
  
FROM employees  
  
START WITH manager_id IS NULL  
  
CONNECT BY PRIOR employee_id = manager_id;
```

○

Miscellaneous

1) How do you create an index on a table in SQL Server?

```
CREATE INDEX index_name  
ON table_name (column_name);
```

2) What are some best practices for designing indexes?

- Analyze query patterns to design effective indexes.
- Limit the number of indexes to avoid performance degradation.
- Use covering indexes to include all columns needed by queries.
- Avoid over-indexing columns with a low cardinality.
- Regularly monitor and maintain indexes to avoid fragmentation.

3) Explain the concept of a composite index and when you might use one.

- A **composite index** is an index on multiple columns of a table. It is used to improve the performance of queries that filter or sort based on the values of these columns.
- **Usage:** When queries frequently filter or sort based on multiple columns together, such as `WHERE column1 = value1 AND column2 = value2`.

Example:

sql

```
CREATE INDEX idx_composite  
ON table_name (column1, column2);
```

Security and Accessibility

1. Imagine you are setting up a new database for an e-commerce website. How would you ensure that only authorized users have access to sensitive customer data?

1. Define Access Control Requirements

- **Identify Sensitive Data:** Determine which data is considered sensitive (e.g., personal information, payment details).

- Determine User Roles: Define roles and responsibilities for users (e.g., administrators, support staff, analysts) and what data each role should access.

2. Implement Database Security Measures

1. Use Database Roles and Permissions

```
CREATE ROLE AdminRole;
```

```
CREATE ROLE SupportRole;
```

```
CREATE ROLE AnalystRole;
```

```
GRANT SELECT, INSERT, UPDATE, DELETE ON Customers TO AdminRole;
```

```
GRANT SELECT ON Customers TO SupportRole;
```

```
GRANT SELECT ON Orders TO AnalystRole;
```

```
ALTER SERVER ROLE AdminRole ADD MEMBER [AdminUser];
```

```
ALTER SERVER ROLE SupportRole ADD MEMBER [SupportUser];
```

```
ALTER SERVER ROLE AnalystRole ADD MEMBER [AnalystUser];
```

2. Use Column-Level Security

- Restrict Access to Sensitive Columns: Limit access to sensitive columns using views or column-level permissions.

```
-- Create a view that excludes sensitive columns
```

```
CREATE VIEW vw_CustomerBasicInfo AS
```

```
SELECT CustomerID, Name, Email
```

```
FROM Customers;
```

```
-- Grant access to the view
```

```
GRANT SELECT ON vw_CustomerBasicInfo TO SupportRole;
```

3. Implement Row-Level Security

- Filter Data Based on User Context: Use row-level security to restrict access to

rows based on user roles or attributes.

```
CREATE FUNCTION dbo.FilterCustomers(@UserRole NVARCHAR(50))
```

```
RETURNS TABLE
```

```
WITH SCHEMABINDING
```

```
AS
```

```
RETURN
```

```
SELECT CustomerID, Name, Email
```

```
FROM Customers
```

```
WHERE
```

```
(@UserRole = 'Admin') OR
```

```
(@UserRole = 'Support' AND Region = 'NorthAmerica'); -- Example condition
```

```
CREATE SECURITY POLICY CustomerSecurityPolicy
```

```
ADD FILTER PREDICATE dbo.FilterCustomers(UserRole) ON dbo.Customers;
```

4. Use Encryption

- **Encrypt Sensitive Data:** Use encryption to protect data at rest and in transit. This includes encrypting sensitive columns and using secure connections (e.g., SSL/TLS).

-- Example of column encryption

```
CREATE MASTER KEY ENCRYPTION BY PASSWORD = 'YourStrongPassword';
```

```
CREATE CERTIFICATE MyCertificate WITH SUBJECT = 'Database Encryption';
```

```
CREATE SYMMETRIC KEY MySymmetricKey WITH ALGORITHM = AES_256 ENCRYPTION  
BY
```

```
CERTIFICATE MyCertificate;
```

```
OPEN SYMMETRIC KEY MySymmetricKey DECRYPTION BY CERTIFICATE MyCertificate;
```

-- Encrypt sensitive column

```
UPDATE Customers SET CreditCardNumber =
```

```
ENCRYPTBYKEY(KEY_GUID('MySymmetricKey'), CreditCardNumber);
```

5. Implement Authentication and Authorization

- Use Strong Authentication Methods: Implement multi-factor authentication (MFA) and strong password policies for accessing the database.
- Enforce Least Privilege Principle: Ensure users have the minimum level of access necessary to perform their duties.

6. Regularly Audit and Monitor Access

- Audit Database Activity: Enable and review auditing to track access and modifications to sensitive data.

-- Example: Enabling SQL Server Audit

```
CREATE SERVER AUDIT Audit_YourDB  
TO FILE ( FILEPATH = 'C:\AuditLogs\' )  
WITH ( QUEUE_DELAY = 1000, ON_FAILURE = CONTINUE );  
CREATE DATABASE AUDIT SPECIFICATION Audit_Spec  
FOR SERVER AUDIT Audit_YourDB  
ADD (SELECT ON dbo.Customers BY [public]) WITH (STATE = ON);
```

MSSQL Install and Configure, OLAP and OLTP

1. A large retail company needs a high-performing OLTP system for processing sales and an OLAP system for analyzing sales data. Explain how you would design and optimize these systems.

OLTP System Design and Optimization:

1. Database Schema Design:

- Normalization: Use normalization techniques to reduce data redundancy and ensure data integrity. Typically, this involves organizing the data into multiple related tables.

- Entity-Relationship Model: Design a schema that captures the business entities (e.g., customers, orders, products) and their relationships. For instance:

- Customers: CustomerID, Name, Address, etc.

- Orders: OrderID, CustomerID, OrderDate, TotalAmount, etc.

- OrderDetails: OrderDetailID, OrderID, ProductID, Quantity, Price, etc.

- Products: ProductID, Name, Category, Price, etc.

2. Indexing:

- Primary Keys: Ensure primary keys are indexed to speed up data retrieval.

- Secondary Indexes: Create secondary indexes on frequently queried fields, such as CustomerID or ProductID, to enhance search performance.

3. Transaction Management:

- ACID Properties: Ensure the system adheres to ACID (Atomicity, Consistency, Isolation, Durability) principles to handle transactions reliably and maintain data integrity.

- Concurrency Control: Implement locking mechanisms or multiversion concurrency control (MVCC) to handle concurrent transactions and prevent issues like dirty reads or lost updates.

4. Performance Optimization:

- Query Optimization: Optimize SQL queries for speed. Use execution plans to understand query performance and adjust indexes or queries accordingly.

- Partitioning: Consider partitioning large tables to improve performance and manageability.

- Caching: Implement caching strategies to reduce database load and improve response times for frequent queries.

5. Scalability:

- Horizontal Scaling: Use database sharding or clustering techniques to distribute the load across multiple servers if necessary.
- Load Balancing: Implement load balancers to evenly distribute database requests and improve availability.

6. Security and Backup:

- Data Security: Implement robust access controls, encryption, and regular security audits to protect sensitive data.
- Backup Strategy: Regularly back up data to prevent loss due to hardware failures or other issues. Consider automated backup solutions.

OLAP System Design and Optimization

1. Database Schema Design:

- Star Schema: Organize data into a central fact table (e.g., Sales) connected to dimension tables (e.g., Time, Product, Customer). This schema is effective for complex queries and reporting.
- Snowflake Schema: A variation where dimension tables are normalized into multiple related tables. This can save storage but might complicate queries.

2. Data Aggregation:

- Pre-Aggregation: Pre-calculate and store aggregate values (e.g., total sales per month) to speed up query responses. Use materialized views or summary tables for this purpose.
- Cube Creation: Design OLAP cubes to aggregate data along various dimensions (e.g., sales by region, time, and product). This facilitates fast multidimensional analysis.

3. Indexing and Storage:

- Bitmap Indexes: Use bitmap indexes on dimension attributes to improve query

performance, especially for categorical data.

- Columnar Storage: Store data in columnar format to improve performance for read-heavy operations and aggregations.

4. Query Optimization:

- Query Optimization: Optimize complex queries for performance. Ensure that they are written to leverage the pre-aggregated data and indexing.
- Execution Plans: Analyze execution plans for queries to identify and resolve bottlenecks.

5. Performance Optimization:

- Data Partitioning: Partition large fact tables by date or other dimensions to improve query performance and manageability.
- Caching: Implement caching mechanisms for frequently accessed data or query results to reduce load on the OLAP system.

6. Scalability:

- Data Warehouse Appliances: Consider using data warehouse appliances or cloud-based solutions that offer scalable storage and compute resources for large-scale analytics.
- Distributed Processing: Use distributed computing frameworks (e.g., Apache Hadoop, Apache Spark) for processing and analyzing large datasets.

7. Data Integration and ETL:

- ETL Processes: Design efficient ETL (Extract, Transform, Load) processes to regularly update the OLAP system with data from the OLTP system. Ensure that these processes are optimized for performance and handle data quality issues.

8. Security and Governance:

- Access Controls: Implement role-based access controls to ensure that only authorized users can access sensitive or critical data.

- Data Governance: Establish data governance policies to ensure data quality, consistency, and compliance with regulations.

Data Encryption and Storage

1. What is authentication vs authorization.

Aspect Authentication Authorization

Definition Verifying the identity of a user or system.

Determining what an authenticated user or system can access or do.

Purpose To confirm that the entity is who they claim to be.

To control access to resources based on permissions or roles.

Process Involves checking credentials (e.g., username/password, biometrics).

Involves checking permissions or roles to grant or deny access.

Key Focus Identity verification. Access control and permissions.

Examples

- Username and password
- Twofactor authentication (2FA)
-

Biometric verification (e.g., fingerprint)

- Role-Based Access Control

(RBAC)
- Access Control Lists

(ACLs)
- Permissions and policies

Occurs First Yes, authentication occurs before authorization. No, it occurs after authentication.

What It Ensures That the entity is legitimate. That the entity has the right permissions.

Implementation

- Login systems
- Security

- tokens
- Identity verification

methods

- Defining roles and permissions
-

Implementing access control policies

Dependency Authorization relies on successful authentication.

Authentication provides the identity needed for authorization.

DDL

1. You need to create a new table with columns for EmployeeID, Name, Position, and Salary.

Write the SQL statement to create this table. Add unique constraint, primary key accordingly

Query:

```
CREATE TABLE Employees ( EmployeeID INT NOT NULL AUTO_INCREMENT, Name  
VARCHAR(100) NOT NULL, Position VARCHAR(50) NOT NULL, Salary DECIMAL(10, 2) NOT  
NULL,  
PRIMARY KEY (EmployeeID), UNIQUE (EmployeeID) );
```

2. A table needs to be altered to add a new column Email, but only if it doesn't already

exist. Write the SQL statement to achieve this.

-- Check if the column exists

IF NOT EXISTS (

SELECT 1

FROM sys.columns

WHERE Name = 'Email'

AND Object_ID = Object_ID('YourTableName')

)

BEGIN

-- Add the column if it does not exist

EXEC sp_executesql N'

ALTER TABLE YourTableName

ADD Email VARCHAR(255);

';

END

3.What is Composite key?

A composite key is a type of primary key in a database that consists of two or more columns to uniquely identify a record in a table. Unlike a single-column primary key, which relies on one column for uniqueness, a composite key uses a combination of multiple columns to ensure that each record is unique.

Key Characteristics of Composite Keys:

1. Uniqueness: The combination of the values in the columns that make up the composite key must be unique for each row in the table. No two rows can have the same combination of values in these columns.

2. Multi-Column: A composite key involves more than one column. The uniqueness is enforced by the combination of all the specified columns.
3. Primary Key Constraint: Composite keys are used as primary keys to ensure that each row in a table can be uniquely identified.
4. Relationships: Composite keys are often used in tables that represent many-to-many relationships between entities or where a single column alone cannot guarantee uniqueness.

Example:

```
CREATE TABLE CourseEnrollments (  
    StudentID INT NOT NULL,  
    CourseID INT NOT NULL,  
    EnrollmentDate DATE,  
    PRIMARY KEY (StudentID, CourseID)  
);
```

DML

1. Write a query to update the Salary column in the Employees table, increasing all salaries by 10%.

Query:

```
UPDATE Employees SET Salary = Salary * 1.10;
```

2. You need to retrieve data from the Orders and Customers tables to find all orders placed by customers from a specific city. Write the query.

Query:

```
SELECT o.OrderID, o.OrderDate, o.TotalAmount, c.CustomerID, c.CustomerName, c.City  
FROM Orders o
```

JOIN Customers c ON o.CustomerID = c.CustomerID

WHERE c.City = 'SpecificCityName';

Aggregate Functions

1. Write a query to find the average Salary in the Employees table, grouped by Department.

Query:

SELECT Department, AVG(Salary) AS AverageSalary

FROM Employees

GROUP BY Department;

2. Describe a scenario where you would use the RANK function to assign ranks to employees based on their sales performance.

WITH SalesTotals AS (

SELECT

e.EmployeeID,

e.Name,

SUM(s.SaleAmount) AS TotalSales

FROM Employees e

JOIN Sales s ON e.EmployeeID = s.EmployeeID

WHERE s.SaleDate BETWEEN '2024-06-01' AND '2024-06-30'

GROUP BY e.EmployeeID, e.Name

),

RankedSales AS (

SELECT

EmployeeID,

Name,

TotalSales,

RANK() OVER (ORDER BY TotalSales DESC) AS SalesRank

FROM SalesTotals

)

SELECT * FROM RankedSales;

Filters

1. Write a query to filter out all products from the Products table that have a Price greater than rs.100.

Query:

```
SELECT * FROM Products WHERE Price > 100;
```

2. Explain how you would use the CASE statement to filter data based on multiple conditions, such as categorizing products into different price ranges.

Example Scenario:

Suppose you have a Products table with a Price column, and you want to categorize products into

price ranges:

- Low: Price less than ₹50
- Medium: Price between ₹50 and ₹100
- High: Price greater than ₹100

Query:

```
SELECT ProductID, ProductName, Price,
```

```
CASE
```

```
WHEN Price < 50 THEN 'Low'
```

```
WHEN Price BETWEEN 50 AND 100 THEN 'Medium'
```

```
WHEN Price > 100 THEN 'High'
```

ELSE 'Unknown'

END AS PriceCategory FROM Products;

Multiple Tables: Normalization

1. A denormalized database is causing performance issues. Describe the steps you would take to normalize it and improve performance till 3NF

1. Identify the Current Schema

- Review the existing schema to understand its structure and relationships.
- Identify the tables and their columns, and look for redundancy and anomalies.

2. Apply First Normal Form (1NF)

- Ensure that each table has a primary key.
- Eliminate repeating groups or arrays by creating separate rows for each instance.
- Ensure that each column contains atomic (indivisible) values.

3. Apply Second Normal Form (2NF)

- Identify and remove partial dependencies: columns should depend on the whole primary key.
- Create separate tables for sets of related data and establish foreign key relationships.
- Ensure that non-key attributes are fully functionally dependent on the entire primary key.

4. Apply Third Normal Form (3NF)

- Remove transitive dependencies: non-key columns should not depend on other non-key columns.
- Create new tables to hold transitive dependencies and link them with foreign keys.
- Ensure that each non-key attribute is non-transitively dependent on the primary

key.

5. Review and Refactor

- Test Performance: Check the performance of the normalized schema with sample queries.
- Optimize Indexing: Add appropriate indexes to improve query performance.
- Update Application Logic: Modify application queries and logic to accommodate the normalized schema.

Indexes & Constraints

1. Write a SQL statement to create an index on the Email column of the Users table.

Query:

```
CREATE INDEX idx_email ON Users (Email);
```

Joins

1. Write a query to perform an inner join between the Orders and Customers tables to retrieve all orders along with customer names.

Query:

```
SELECT
```

```
o.OrderID,
```

```
o.OrderDate,
```

```
o.TotalAmount,
```

```
c.CustomerID,
```

```
c.CustomerName
```

```
FROM Orders o
```

```
INNER JOIN Customers c ON o.CustomerID = c.CustomerID;
```

2. Describe a scenario where a full outer join would be necessary, and provide a query

example.

Scenario: Employee Department Mismatches

Suppose you have two tables: Employees and Departments.

- Employees Table: Contains information about employees, including their department assignments.
- Departments Table: Contains information about departments, including their status.

You want to generate a report showing all employees and all departments, even if some employees are not assigned to any department or some departments do not have any employees.

Query:

```
SELECT
e.EmployeeID,
e.EmployeeName,
e.DepartmentID AS EmployeeDepartmentID,
d.DepartmentName,
d.DepartmentID AS DepartmentID
FROM Employees e
FULL OUTER JOIN Departments d ON e.DepartmentID = d.DepartmentID;
```

Alias

1. Write a query using table aliases to simplify a complex join between three tables: Orders, Customers, and Products.

Query:

```
SELECT
o.OrderID AS OrderID,
o.OrderDate AS OrderDate,
```

```
o.Quantity AS Quantity,  
c.CustomerName AS CustomerName,  
p.ProductName AS ProductName  
FROM Orders o  
INNER JOIN Customers c ON o.CustomerID = c.CustomerID  
INNER JOIN Products p ON o.ProductID = p.ProductID;
```

2.Explain how column aliases can be used in a query with aggregate functions to make the results more readable.

Query:

```
SELECT  
ProductID,  
SUM(Quantity) AS TotalQuantity,  
SUM(SaleAmount) AS TotalSales,  
AVG(SaleAmount) AS AverageSaleAmount  
FROM Sales  
GROUP BY ProductID;
```

Joins vs Sub Queries

1.Provide a scenario where a subquery would be more appropriate than a join, and write the corresponding query.

Scenario:

Suppose you have an e-commerce database with two tables: Orders and Customers.

- Orders: Contains details about customer orders.
- order_id (Primary Key)
- customer_id

- order_amount
- order_date
- Customers: Contains details about customers.
- customer_id (Primary Key)
- customer_name
- customer_email

You want to find customers who have placed more than 5 orders. In this scenario, a subquery can be more appropriate than a join because you need to filter based on an aggregated result (the number of orders per customer) rather than combining rows from both tables.

Query:

```
SELECT customer_id, customer_name, customer_email
FROM Customers
WHERE customer_id IN (
  SELECT customer_id
  FROM Orders
  GROUP BY customer_id
  HAVING COUNT(order_id) > 5
);
```

2. Compare the performance implications of using a join versus a subquery for retrieving data from large tables.

Joins

Pros:

- Efficient with Indexes: Fast if proper indexes are in place.
- Optimized by DBMS: Modern systems optimize joins well.

Cons:

- Complexity and Cost: Can be expensive with large datasets and no indexes.
- Intermediate Results: May involve costly intermediate data handling.

Subqueries

Pros:

- Granular Filtering: Useful for complex filters and aggregated results.
- Avoids Cartesian Products: Can prevent unnecessary large intermediate results.

Cons:

- Execution Overhead: May be less efficient if the DBMS executes the subquery repeatedly.
- Less Predictable: Performance can vary based on how the subquery is optimized.

Types

1. Describe a scenario where using a DATETIME data type would be essential, and explain how you would store and retrieve this data.

Storing Data

Table Schema:

```
CREATE TABLE UserActivity (  
    activity_id INT AUTO_INCREMENT PRIMARY KEY,  
    user_id INT NOT NULL,  
    login_time DATETIME NOT NULL,  
    logout_time DATETIME,  
    FOREIGN KEY (user_id) REFERENCES Users(user_id)  
);
```

Example Insert Statement:

```
INSERT INTO UserActivity (user_id, login_time, logout_time)
VALUES (1, '2024-07-27 08:30:00', '2024-07-27 10:15:00');
```

Retrieving Data

```
SELECT user_id, login_time, logout_time
FROM UserActivity WHERE DATE(login_time) = '2024-07-27';
```

Query to Calculate Session Durations:

```
SELECT user_id, login_time, logout_time,
TIMESTAMPDIFF(MINUTE, login_time, logout_time) AS session_duration_minutes
FROM UserActivity
WHERE user_id = 1;
```

2.Explain how you would handle data type conversions when importing data from a CSV file with mixed data types.

- Understand the CSV File: Know the structure and data types.
- Prepare the Schema: Define the database schema with appropriate data types.
- Pre-process and Convert Data: Clean and convert data formats before importing.
- Handle Errors: Log and correct conversion issues.
- Verify and Validate: Ensure data is correctly imported and valid.

Correlation and Non-Correlation

1.Write a query using a correlated subquery to find all employees who have a salary higher than the average salary in their department.

Query:

```
SELECT e1.employee_id, e1.employee_name, e1.salary, e1.department_id
FROM Employees e1
WHERE e1.salary > (
```

```
SELECT AVG(e2.salary)
FROM Employees e2
WHERE e2.department_id = e1.department_id
);
```

2.Explain how non-correlated subqueries can be optimized for better performance compared to correlated subqueries.

Non-Correlated Subqueries

Definition:

- A non-correlated subquery does not reference any columns from the outer query. It is executed once, and its result is used by the outer query.

Correlated Subqueries

Definition:

- A correlated subquery references columns from the outer query and is executed for each row of the outer query. This means it is executed multiple times, once for each row in the outer query.

Key Differences and Performance Considerations:

1. Execution Frequency:

- Non-Correlated Subqueries: Executed once and reused.
- Correlated Subqueries: Executed multiple times, once for each row in the outer query.

2. Optimization Potential:

- Non-Correlated Subqueries: Often more straightforward to optimize due to single execution and result caching.
- Correlated Subqueries: Requires techniques to minimize repeated

execution, such as rewriting to joins or using temporary tables.

3. Use Cases:

- Non-Correlated Subqueries: Best for operations where the subquery provides a static value or set of values used by the outer query.
- Correlated Subqueries: Useful for operations where the result depends on the row being processed by the outer query but can often be optimized by rethinking the query structure.

Introduction to TSQL, Procedures, Functions, Triggers, Indices

1. Write a simple stored procedure to insert a new record into the Employees table.

Example Table Structure

```
CREATE TABLE Employees (  
    employee_id INT AUTO_INCREMENT PRIMARY KEY,  
    employee_name VARCHAR(100) NOT NULL,  
    department_id INT NOT NULL,  
    salary DECIMAL(10, 2) NOT NULL,  
    hire_date DATE NOT NULL  
);
```

Stored Procedure Definition

```
DELIMITER //  
  
CREATE PROCEDURE InsertEmployee(  
    IN p_employee_name VARCHAR(100),  
    IN p_department_id INT,  
    IN p_salary DECIMAL(10, 2),  
    IN p_hire_date DATE
```

)

BEGIN

INSERT INTO Employees (employee_name, department_id, salary, hire_date)

VALUES (p_employee_name, p_department_id, p_salary, p_hire_date);

END //

DELIMITER ;

How to Call the Stored Procedure

CALL InsertEmployee('John Doe', 2, 55000.00, '2024-07-27');

2. Describe a scenario where you would use a trigger to enforce business rules.

Scenario: Enforcing Minimum Salary

Business Rule: Employees must have a salary of at least \$30,000.

Using a Trigger:

1. Table Definition:

```
CREATE TABLE Employees (  
    employee_id INT AUTO_INCREMENT PRIMARY KEY,  
    employee_name VARCHAR(100) NOT NULL,  
    department_id INT NOT NULL,  
    salary DECIMAL(10, 2) NOT NULL,  
    hire_date DATE NOT NULL  
);
```

2. Create Triggers:

DELIMITER //

CREATE TRIGGER CheckSalaryBeforeInsert

BEFORE INSERT ON Employees

```

FOR EACH ROW
BEGIN
    IF NEW.salary < 30000.00 THEN
        SIGNAL SQLSTATE '45000'
        SET MESSAGE_TEXT = 'Salary cannot be less than $30,000.';
    END IF;
END //

CREATE TRIGGER CheckSalaryBeforeUpdate
BEFORE UPDATE ON Employees
FOR EACH ROW
BEGIN
    IF NEW.salary < 30000.00 THEN
        SIGNAL SQLSTATE '45000'
        SET MESSAGE_TEXT = 'Salary cannot be less than $30,000.';
    END IF;
END //

DELIMITER ;

```

Security and Accessibility

1. Mention any 2 of the common security measures to protect a SQL Server database?

- Authentication and Authorization:

This involves configuring who can access the SQL Server and what they can do once they have access. SQL Server supports both Windows Authentication (using Windows accounts) and SQL Server Authentication (using SQL Server-specific accounts).

Implementing strong, unique passwords and using roles and permissions to grant the least

privileges necessary to users helps minimize security risks.

- Encryption:

This involves protecting data both at rest and in transit. For data at rest, SQL Server provides features like Transparent Data Encryption (TDE) to encrypt the database files. For data in transit, you can use SSL/TLS to encrypt the data as it travels between the client and the server, ensuring that sensitive information isn't exposed during transmission.

2. How do you create a user and assign roles in MSSQL?

1. Creating a Login

```
CREATE LOGIN [NewLogin] WITH PASSWORD = 'YourStrongPassword';
```

2. Creating a Database User

```
CREATE USER [NewUser] FOR LOGIN [NewLogin];
```

3. Assigning Roles to the User

```
ALTER ROLE db_datareader ADD MEMBER [NewUser];
```

```
ALTER ROLE db_datawriter ADD MEMBER [NewUser];
```

3. Explain how encryption can be implemented for data at rest in MSSQL.

1. Transparent Data Encryption (TDE)

TDE encrypts the entire database, including the data files, log files, and backups. This is done transparently without requiring changes to the application or the database schema. Here's how you can implement TDE:

1. Create a Master Key:

```
CREATE MASTER KEY ENCRYPTION BY PASSWORD =  
'YourStrongPassword';
```

2. Create a Certificate:

```
CREATE CERTIFICATE MyTDECert
```


WITH SUBJECT = 'TDE Certificate';

3. Create a Database Encryption Key (DEK):

CREATE DATABASE ENCRYPTION KEY

WITH ALGORITHM = AES_256 ENCRYPTION BY SERVER CERTIFICATE MyTDECert;

4. Enable Encryption on the Database:

ALTER DATABASE YourDatabaseName

SET ENCRYPTION ON;

5. Backup the Certificate and Master Key:

BACKUP CERTIFICATE MyTDECert

TO FILE = 'C:\Backup\MyTDECert.cer'

WITH PRIVATE KEY (

FILE = 'C:\Backup\MyTDECert.pvk',

ENCRYPTION BY PASSWORD = 'YourStrongPassword'

);

BACKUP MASTER KEY

TO FILE = 'C:\Backup\MasterKey.bak'

ENCRYPTION BY PASSWORD = 'YourStrongPassword';